

```
In [1]: # import relevant libraries
import numpy as np
import matplotlib.pyplot as plt
import keras.layers as l
import keras.optimizers as o
import keras.models as m
import keras
from keras.utils import to_categorical
```

Using TensorFlow backend.


```
In [3]: # load cifar10 dataset
from keras.datasets import cifar10

(x_train,y_train), (x_test,y_test) = cifar10.load_data()
# rescale pixel values
x_train = (1/255.0)*x_train
x_test = (1/255.0)*x_test
# labels -> one-hot encodings
y_train = to_categorical(y_train.reshape([-1, 1]))
y_test = to_categorical(y_test.reshape([-1, 1]))
```

```
In [4]: # for jupyter notebook (ignore)
%matplotlib inline
```

```
In [5]: # plot image
plt.imshow(x_train[0])
```

Out[5]: <matplotlib.image.AxesImage at 0x7ff7359c2ef0>



## Sequential API

```
In [6]: # define model using sequential api
model = m.Sequential([
    l.Conv2D(32, (3,3), padding='same', activation='relu', input_shape=(32,
    l.MaxPooling2D((3,3)),
    l.Conv2D(64, (3,3), padding='same', activation='relu'),
    l.MaxPooling2D((3,3)),
    l.Flatten(),
    l.Dense(32, activation='relu'),
    l.Dense(10, activation="softmax")
])

# print a model summary
model.summary()
```

WARNING: Logging before flag parsing goes to stderr.

W1230 13:55:03.580109 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:74: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

W1230 13:55:03.607579 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

W1230 13:55:03.611763 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:4138: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

W1230 13:55:03.630030 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:3976: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 10, 10, 32)	0
conv2d_2 (Conv2D)	(None, 10, 10, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 64)	0
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 32)	18464
dense_2 (Dense)	(None, 10)	330
=====		
Total params: 38,186		
Trainable params: 38,186		
Non-trainable params: 0		
=====		

In [7]: `# compile a model`

```
model.compile(optimizer=o.Adam(),
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

W1230 13:55:03.697125 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

W1230 13:55:03.705946 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:3295: The name tf.log is deprecated. Please use tf.math.log instead.

In [8]: `# fit model to data`

```
model.fit(x_train, y_train, epochs=10, validation_data = (x_test, y_test))
```

W1230 13:55:03.814929 140701437798144 deprecation.py:323] From /home/gpik/miniconda3/lib/python3.7/site-packages/tensorflow/python/ops/math\_grad.py:1250: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

W1230 13:55:03.862976 140701437798144 deprecation\_wrapper.py:119] From /home/gpik/miniconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:986: The name tf.assign\_add is deprecated. Please use tf.compat.v1.assign\_add instead.

Train on 50000 samples, validate on 10000 samples

Epoch 1/10

50000/50000 [=====] - 15s 306us/step - loss: 1.5707 - acc: 0.4331 - val\_loss: 1.2921 - val\_acc: 0.5419

Epoch 2/10

50000/50000 [=====] - 14s 290us/step - loss: 1.1855 - acc: 0.5826 - val\_loss: 1.1453 - val\_acc: 0.5940


Epoch 3/10

## Functional API

```
In [9]: from keras.datasets import mnist

# load mnist data
(x_train,y_train), (x_test,y_test) = mnist.load_data()
# rescale pixel values
x_train = (1/255.0)*x_train
x_test = (1/255.0)*x_test
# labels -> one-hot encodings
y_train = to_categorical(y_train.reshape([-1, 1]))
y_test = to_categorical(y_test.reshape([-1, 1]))
plt.imshow(x_train[0])
```

Out[9]: <matplotlib.image.AxesImage at 0x7ff73471e898>



```
In [10]: # define model & compile
input_layer = l.Input((28,28))
Intermediate_layer = l.Flatten()(input_layer)
Intermediate_layer=l.Dense(256, activation='relu')(Intermediate_layer)
output_layer = l.Dense(10, activation="softmax")(Intermediate_layer)

model = m.Model(input_layer, output_layer)
model.summary()
model.compile(optimizer=o.SGD(),
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 28, 28)	0
flatten_2 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 256)	200960
dense_4 (Dense)	(None, 10)	2570
Total params: 203,530		
Trainable params: 203,530		
Non-trainable params: 0		

```
In [12]: model.fit(x_train.v_train, epochs=10, validation_data = (x_test.v_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10  
60000/60000 [=====] - 3s 55us/step - loss: 0.6223  
- acc: 0.8451 - val\_loss: 0.3476 - val\_acc: 0.9071

Epoch 2/10  
60000/60000 [=====] - 3s 53us/step - loss: 0.3284  
- acc: 0.9082 - val\_loss: 0.2839 - val\_acc: 0.9204

Epoch 3/10  
60000/60000 [=====] - 3s 53us/step - loss: 0.2797  
- acc: 0.9218 - val\_loss: 0.2522 - val\_acc: 0.9309

Epoch 4/10  
60000/60000 [=====] - 3s 53us/step - loss: 0.2490  
- acc: 0.9309 - val\_loss: 0.2309 - val\_acc: 0.9333

Epoch 5/10  
60000/60000 [=====] - 3s 52us/step - loss: 0.2260  
- acc: 0.9372 - val\_loss: 0.2110 - val\_acc: 0.9391

Epoch 6/10  
60000/60000 [=====] - 3s 52us/step - loss: 0.2072  
- acc: 0.9423 - val\_loss: 0.1973 - val\_acc: 0.9444

Epoch 7/10  
60000/60000 [=====] - 3s 52us/step - loss: 0.1916  
- acc: 0.9468 - val\_loss: 0.1832 - val\_acc: 0.9482

Epoch 8/10  
60000/60000 [=====] - 3s 53us/step - loss: 0.1781  
- acc: 0.9502 - val\_loss: 0.1716 - val\_acc: 0.9509

Epoch 9/10  
60000/60000 [=====] - 3s 52us/step - loss: 0.1665  
- acc: 0.9541 - val\_loss: 0.1624 - val\_acc: 0.9533

Epoch 10/10  
60000/60000 [=====] - 3s 55us/step - loss: 0.1563  
- acc: 0.9564 - val\_loss: 0.1558 - val\_acc: 0.9551

```
Out[12]: <keras.callbacks.History at 0x7ff7347d0c18>
```

### Obtain model gradients

```
In [1]: import keras.backend as K

# function to obtain grads for each parameter
def get_gradients(model, inputs, outputs):
    grads = model.optimizer.get_gradients(model.total_loss, model.trainable_variables)
    symb_inputs = (model._feed_inputs + model._feed_targets + model._feed_sample_weights)
    f = K.function(symb_inputs, grads)
    x, y, weight = model._standardize_user_data(inputs, outputs)
    output_grad = f(x + y + weight)
    return np.array(output_grad)
```

Using TensorFlow backend.

```
In [23]: # test
         grads = get_gradients(model, x_train[0:1], y_train[0:1])
```

```
In [33]: # Print max from each layer, weight and biases. Even indices hold biases
print(grads.shape)
for i, _ in enumerate(grads):
    print(grads[i].shape)
    max_gradient_layer_i = np.max(grads[i])
    print(max_gradient_layer_i)


(4,)
(784, 256)
0.0068208314
(256,)
0.0068208314
(256, 10)
0.034386244
(10,)
0.011868487
```

### Load .mat files

```
In [50]: from scipy.io import loadmat

SVHN_directory = "train_32x32.mat"
# load .mat file
data_raw = loadmat(SVHN_directory)
data = np.array(data_raw['X'])
# make correct shape
data = np.moveaxis(data, -1, 0)
print(data.shape)
plt.imshow(data[0])
labels = data_raw['y']
print(labels.shape)
print(labels[0])

(73257, 32, 32, 3)
(73257, 1)
[1]
```




### Make your own digits

```
In [96]: import cv2 as cv
from skimage.transform import resize

image_dir = "test_digit2.jpg"
# load & smoothen image
kernel = np.ones((7,7),np.float32)/49
image = cv.imread(image_dir,cv.IMREAD_GRAYSCALE)
image = cv.filter2D(image,-1,kernel)

# make numpy array
image = np.array(image)
image = resize(image, (28,28))
# make negative
image = np.ones(image.shape) - image

plt.imshow(image, cmap="gray")
plt.show()
```



```
In [97]: print(image.shape)
# predict label
np.argmax(model.predict(np.array([image])))
(28, 28)
```

Out[97]: 8

In [ ]: